



Give Your Personal AI Financial Analyst the Ability to Reason

A Head-to-Head Benchmark

10 questions across 5 categories. 5 companies: AAPL, MSFT, GOOGL, AMZN, META.
Scored across Accuracy, Depth, Cross-Document Synthesis, Actionability, and Specificity.

CoWork standalone vs. CoWork + Vrin MCP


You are receiving this report because we admire your work and believe your audience would benefit from what Vrin can do. We would love for you to try it, and if you find it valuable, we would be thrilled to explore a partnership together.

vrin.cloud

Knowledge Reasoning for Financial Intelligence

Executive Summary

We gave two systems the same task: answer 10 financial analysis questions about Apple, Microsoft, Google, Amazon, and Meta using their real SEC filings (annual reports, earnings releases, and earnings call transcripts). One system was CoWork standalone, which had access to these filings locally along with web search. The other was CoWork + Vrin MCP, which added Vrin's knowledge reasoning engine on top. Same AI model, same questions, scored blind across five dimensions.



Three Findings That Matter

MULTI-SOURCE REASONING

When asked how much Meta has actually committed to AI infrastructure, Vrin pulled three separate disclosure sources from the annual report: \$72.2B in capital expenditures, \$81.2B in non-cancelable purchase commitments, and \$58.1B in lease obligations. Together, that is \$209B in total committed spending. The standalone system found only the headline capex figure, understating Meta's true exposure by roughly two-thirds.

SELF-CORRECTION

Vrin caught and corrected its own earlier mistake. In the first question, it reported that Meta's free cash flow grew 19.5%. Two questions later, while analyzing capital spending trajectories, it re-read the actual filing, realized the figure was wrong, and corrected it to a 16.3% decline (\$43.6B). No human prompted this. The system cross-checked itself against source data.

CROSS-FILING PATTERN DETECTION

When asked to spot differences between what companies say on earnings night versus what they disclose in annual reports, Vrin found that Microsoft's quarterly OpenAI investment swings wildly: a \$3.1B loss one quarter, a \$7.6B gain the next. These swings reverse the direction of Microsoft's earnings gap each quarter, a pattern only visible when reading across multiple filings.

The floor effect: Vrin raised the minimum per-question score from 74 to 89. It does not just improve the best answers. It eliminates the worst ones.

The Bottom Line

Vrin won 9 of 10 questions. The sole exception (Q6) rewarded creative strategic framing, the type of analysis where AI already performs well without structured document retrieval. For every question that required pulling precise data from filings, cross-referencing disclosures, or verifying claims against source documents, Vrin outperformed consistently.

Methodology: Five dimensions scored 0-20 each (100 per question, 1,000 max). Full responses: github.com/Vrin-cloud/vrin-benchmarks/tree/main/financial-analyst

Category-by-Category Breakdown

Q#	Category / Question	Standalone	Vrin MCP	Delta
Q1	Revenue, margins, FCF comparison	80	89	+9
Q2	AI/cloud capex trajectories	74	93	+19
Q3	META 10-K new risk factors	75	95	+20
Q4	Amazon cash flow breakdown	84	93	+9
Q5	Declining margins + rising R&D;	83	92	+9
Q6	Most aggressive AI capex bet	91	89	-2
Q7	Portfolio revenue concentration	86	91	+5
Q8	Revenue growth guidance ranking	85	91	+6
Q9	8-K vs 10-K inconsistencies	82	95	+13
Q10	GAAP vs non-GAAP earnings gap	84	94	+10
TOTAL		824	922	+98

Cross-Company Comparison (Q1, Q2)

The first two questions asked both systems to compare all five companies side by side. Q1 focused on the fundamentals: revenue growth, profit margins, and free cash flow. Q2 went deeper into how much each company is spending on AI and cloud infrastructure.

Both systems produced solid comparative tables. But Vrin's responses were measurably more precise: figures quoted to the exact million from SEC filings rather than rounded to the nearest billion from web sources. More importantly, Vrin cross-referenced data across questions.

Q2: AI/Cloud Capex Trajectories (+19, Largest Delta)

This question asked how much each company is investing in AI infrastructure and where those trajectories are headed. The standalone system delivered a strong overview with correct directional analysis. Vrin went further in three ways. First, it pulled Meta's exact capital expenditure from the 10-K (\$72.2B including finance leases), whereas standalone reported \$69.7B from a web source. Second, it retrieved verbatim quotes from CFO

earnings calls, including Amy Hood's breakdown that "roughly two-thirds" of Microsoft's capex went to short-lived GPU and CPU assets. Third, it opened its response by correcting a Meta free cash flow figure from the prior question after re-reading the actual 10-K. That self-correction, unprompted by any human, is only possible when the system can reason across its own earlier outputs and source documents simultaneously.

Deep-Dive Filing Analysis (Q3, Q4)

These questions required going deep into a single company's filings. Q3 asked: what are the new risk factors in Meta's latest annual report that were not prominent before? Q4 asked: why is Amazon's operating cash flow so much higher than its net income, and where is all that cash going?

Q3: META's New 10-K Risk Factors (+20)

The standalone system identified three risk categories (AI overbuilding, energy commitments, open-source AI liability) and provided strong narrative analysis, but explicitly acknowledged it "was unable to access the full SEC filing text directly" and relied on third-party news sources instead. Vrin had no such limitation. It pulled three verbatim quotes from Meta's 10-K risk factor section, cross-referenced three separate disclosure tables to calculate \$209B in total infrastructure obligations (\$72.2B capex + \$81.2B contracts + \$58.1B leases), and identified the useful-life extension as a "double-edged sword" that reduces near-term depreciation but increases impairment risk. That level of filing-level synthesis was simply not possible without structured document access.

Q4: Amazon Cash Flow (+9). Both systems produced strong waterfall analyses explaining how Amazon turns \$77.7B in net income into \$139.5B in operating cash flow. Vrin added two filing-specific details: a server useful-life accounting change that added \$889M to depreciation expense, and a \$2.5B FTC settlement accrual. These one-time items, buried in the footnotes, are exactly the kind of details that separate a surface-level summary from a thorough filing review.

Multi-Hop Reasoning (Q5, Q6)

These questions tested the ability to connect information from different parts of multiple filings. Q5 asked which company has both declining profit margins and rising R&D; spending. Q6 asked which company is making the most aggressive AI infrastructure bet relative to its revenue.

Q5 (+9): Both systems correctly identified Microsoft as the answer. But Vrin surfaced a non-obvious insight it called "hidden R&D.;" Microsoft's AI infrastructure costs flow through cost of revenue (which hurts gross margins) rather than the R&D; line. When you combine both, Microsoft's total "innovation spending" actually rose from 42.3% to 42.7% of revenue. Vrin backed this up with the exact 10-K language: "Gross margin percentage decreased slightly driven by continued investments in AI infrastructure." That reframing changes the investment thesis from "margins are deteriorating" to "the company is investing through a different line item."

Q6 (-2, Standalone Wins): Asked which company's AI bet is the most aggressive, the standalone system delivered the stronger answer. It created a memorable "Infrastructure Maximalists vs. Integration Minimalists" framework with one-word strategic labels for each company. Vrin provided precise data (Meta at 57-67% of revenue) but lacked the same narrative punch. This was the only question where the standalone system outscored Vrin.

The pattern here is clear: when the question rewards creative framing and strategic synthesis, standalone competes effectively. When it demands filing-level precision, Vrin pulls ahead.

Portfolio and Correlation (Q7, Q8)

These questions took a portfolio view. Q7 asked: if you own equal amounts of all five stocks, what is your actual business exposure by category? Q8 asked: which management teams are guiding for the strongest growth next year?

Q7 (+5): Both systems produced useful portfolio decompositions. Vrin identified that digital advertising represents 37% of the combined portfolio, the single largest concentration. It also flagged a "hidden dependency" between Apple and Google: the \$20-26B annual payment Google makes to be Safari's default search engine ties the two companies together in ways that are not obvious from their revenue segments alone.

Q8 (+6): Both systems ranked the five companies by forward guidance. Vrin's edge was anchoring each ranking to verbatim management quotes from earnings calls rather than paraphrased analyst summaries. For example, it cited Amazon CFO Brian Olsavsky's hedged language about macro uncertainty alongside Amazon's \$200B capex commitment, noting the tension: conservative revenue guidance paired with aggressive spending signals management expects 18-20%+ demand growth internally.

Anomaly and Inconsistency Detection (Q9, Q10)

These questions tested something subtle: the ability to notice when companies say one thing in their earnings press release and something different in their annual report. This is where structured document reasoning created the widest qualitative gap between the two systems.

Q9: 8-K vs 10-K Inconsistencies (+13)

Both systems identified meaningful gaps across all five companies. The standalone system produced strong narrative analysis of how Amazon de-emphasized its 71% free cash flow decline in the press release, and how Google's \$26B in annual search default payments face an existential antitrust threat. But Vrin went deeper with filing-level evidence: Microsoft's quarterly OpenAI investment swings (\$523M loss, then \$939M loss, then \$3.1B loss, then a \$7.6B gain), Apple's \$1.4B tariff cost disclosed on the call but absent from the press release, and Eddie Cue's court testimony revealing Safari search volumes declined for the first time in 22 years. Vrin also quoted specific 10-K risk disclosure language verbatim, while the standalone system paraphrased without citing the actual filing text.

Q10: GAAP vs Non-GAAP Earnings (+10). Both systems identified the largest gaps between reported and adjusted earnings. Vrin's structural insight was that Microsoft's gap reverses direction every quarter because of OpenAI equity-method accounting: one quarter shows a huge loss, the next shows a huge gain. It also quantified the combined stock-based compensation across all five companies at over \$105B annually, framing it through Warren Buffett's lens as "the most egregious example of accounting that is out of touch with economic reality."

The Pattern: When Does Vrin MCP Matter Most?

Vrin MCP does not replace analytical judgment. It gives your AI assistant the ability to reason across documents with the same rigor you would apply manually.

Dimension	Vrin MCP Wins	Standalone Wins
Question Type	Filing analysis, cross-document retrieval, anomaly detection	Forward-looking strategy synthesis, creative framing
Data Source	Direct 10-K/10-Q quotes, earnings call transcripts, structured tables	Web articles, analyst commentary, synthesized narrative
Win Rate	9 of 10 questions	1 of 10 questions
Score Range	Floor: 89 / Ceiling: 95	Floor: 74 / Ceiling: 91
Unique Capability	Self-correction, filing-level quotes, cross-disclosure synthesis	Memorable frameworks, creative metaphors

Why the Score Gap Understates the Impact

The aggregate score captures quantity. The table below captures quality: Vrin wins on the dimensions that matter most for professional research.

Capability	Vrin MCP	Standalone
Direct SEC filing quotes per response	3-8 verbatim	0
Management earnings call quotes	5-10 verbatim	0-2 paraphrased
Source citations per response	6-12	2-6
Self-corrections against source data	Yes (Q2)	None
Admitted filing access failure	Never	Yes (Q3)
Cross-disclosure discoveries	Multiple (Q3, Q5, Q9, Q10)	None

Just like how you give CoWork access to a local folder on your device, you upload your documents to Vrin once and it builds a knowledge reasoning layer on top. Every time your AI assistant needs to answer a question, it reasons through the actual source material rather than relying on summaries or web search results. Vrin does not replace your AI. It makes every answer filing-grade.



Try Vrin MCP free. Connect it to your AI assistant in minutes.

See the difference filing-level precision makes.

vrin.cloud

Knowledge Reasoning for Financial Intelligence